



Elasticsearch: A Game Changer

Guillaume Delporte, Aly Abdelaleem, Jordi Hoorelbeke

24/05/2023



Objective

Gain a **comprehensive understanding of Elasticsearch**, including its **advanced capabilities**, and discover **how leading companies leverage its powerful features** for various applications.

What We Cover.

1. An **Introduction** to Elasticsearch.
2. Deep Dive into Elasticsearch's **Architecture**.
3. Elasticsearch in Action: **How Companies Utilize It**.
4. **Advanced Search Techniques** with Elasticsearch.
5. **Kibana** and Elasticsearch: The Perfect Combination.
6. **Conclusion**: Key Takeaways on Elasticsearch.

What We Cover.

1. An **Introduction** to Elasticsearch.
2. Deep Dive into Elasticsearch's **Architecture**.
3. Elasticsearch in Action: **How Companies Utilize It**.
4. **Advanced Search Techniques** with Elasticsearch.
5. **Kibana** and Elasticsearch: The Perfect Combination.
6. **Conclusion**: Key Takeaways on Elasticsearch.

An **Introduction** to Elasticsearch.

What is  **elasticsearch**?

- ❑ Elasticsearch: Distributed, Open-Source search, and analytics engine.
- ❑ Handles various types of data: textual, numerical, structured, unstructured.
- ❑ Designed for scalability and can search and index diverse document formats.
- ❑ Created by Shay Banon between 2004-2009



Shay Banon, Elasticsearch's creator



elasticsearch and

APACHE LUCENE™

- ❑ Based on Apache Lucene.
- ❑ Lucene: Free and open-source information retrieval software library.
- ❑ Lucene developed by Doug Cutting.
- ❑ Elasticsearch leverages Lucene's capabilities for full-text search.



Doug Cutting,
Open-Source software developer




Figure from : <https://www.endava.com/en/blog/Engineering/2021/Elasticsearch-and-apache-lucene-fundamentals-behind-the-relevance-score>

How Does elasticsearch Work?

- ❑ Data stored as JSON (JavaScript Object Notation).
- ❑ Document organized in indices, similar as in relational databases.
- ❑ Indices further divided into shards, distributed across cluster nodes.


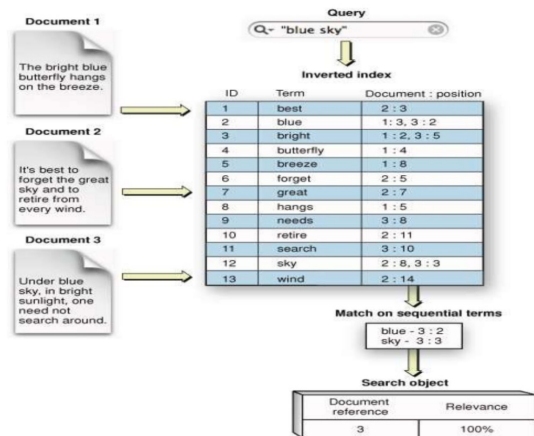



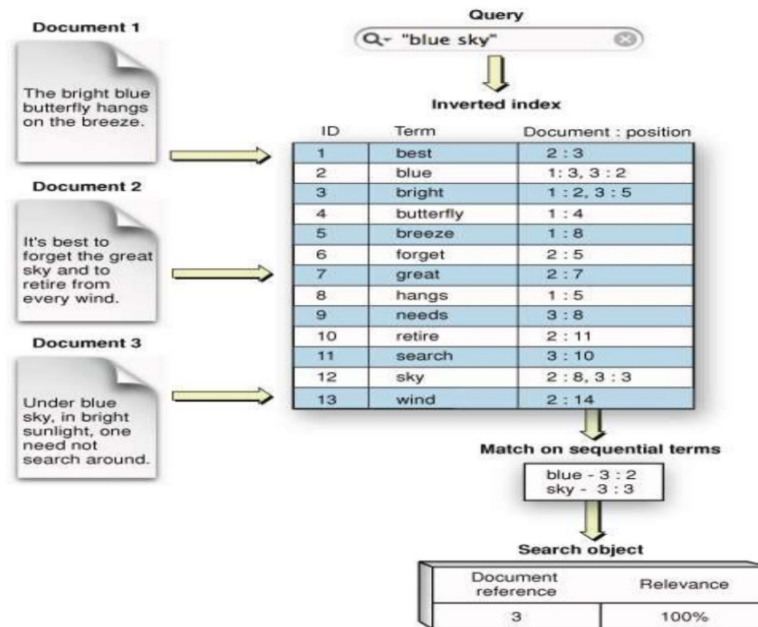
Figure from : <https://codingexplained.com/coding/elasticsearch/understanding-sharding-in-elasticsearch>



Credit : https://developer.apple.com/library/macos/documentation/userexperience/conceptual/SearchKitConcepts/searchKit_basics/searchKit_basics.html

Inverted Index for Speed and Flexibility

- ❏  **elasticsearch** uses an inverted index for fast search.
- ❏ Inverted index stores unique words and documents they appear in.
- ❏ Enables quick lookup and retrieval of search results.



Credit: https://developer.apple.com/library/mac/documentation/userexperience/conceptual/SearchKitConcepts/searchKit_basics/searchKit_basics.html

elasticsearch and the **elastic stack**

- ❑ Elasticsearch is part of the Elastic Stack.
- ❑ Elastic Stack components: Elasticsearch, Logstash, Kibana, Beats.
- ❑ Logstash: Data processing pipeline for ingesting and transforming data.
- ❑ Kibana: Visualization and analysis layer for Elasticsearch data.

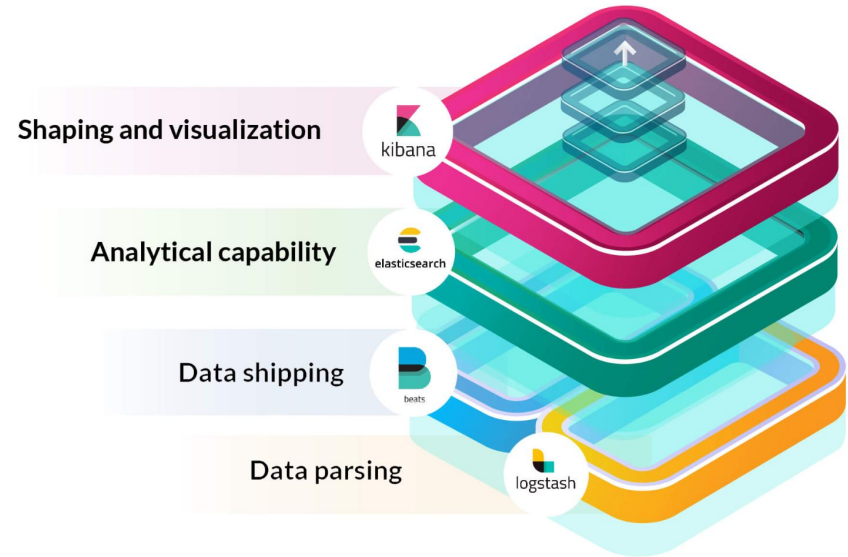



Figure from: <https://quintagroup.com/services/the-elastic-stack-and-its-components-elasticsearch-kibana-logstash-and-beats>

The elastic stack (Cont'd.)

- ❑ Beats: Lightweight data shippers for sending operational data.
- ❑ Beats can send data directly to Elasticsearch or via Logstash for processing.
- ❑ The Elastic Stack provides a comprehensive data management and analysis solution.




What We Cover.

1. An **Introduction** to Elasticsearch.
2. Deep Dive into Elasticsearch's **Architecture**.
3. Elasticsearch in Action: **How Companies Utilize It**.
4. **Advanced Search Techniques** with Elasticsearch.
5. **Kibana** and Elasticsearch: The Perfect Combination.
6. **Conclusion**: Key Takeaways on Elasticsearch.


Deep Dive into **elasticsearch** Architecture

- ❏ A cluster contains one or more nodes.
- ❏ A node is an instance of Elasticsearch
- ❏ Index is used to group shards that contain related data.
- ❏ Shard is where data is stored.
- ❏ Data is stored as documents.



Cluster


- ❑ A cluster is a collection of nodes that work together to achieve a common goal.
- ❑ Each cluster has a unique name.
- ❑ When creating a node, one cluster comes with it by default.
- ❑ The cluster is managed by a master node.



Node


- ❑ A node is an instance of Elasticsearch with a unique id and a name.
- ❑ A node must belong to a cluster.
- ❑ Nodes are recommended to be distributed across several machines, but you can have multiple nodes on one machine.
- ❑ A node can have multiple roles.

Cluster




Index

- ❏ Index is a logical way to group shards that contain documents that are related to each other.
- ❏ Similar to a database in traditional systems.
- ❏ When creating an index, one shard comes with it by default.
- ❏ Index can be configured to allow multiple shards across multiple nodes (Sharding).



Shard

- ❏ Shard is actually where data is stored on disk and where searches are run.
- ❏ A shard can be:
 - ❏ Primary shard
 - ❏ Replica shard
- ❏ Elasticsearch allows you to configure the number of replica shards per primary shard.
- ❏ Benefits of sharding: horizontal scaling, parallel queries, backup, and load balancing using replica shards.



Document

- ❑ Data is stored as documents, which is a JSON object that is stored with a unique id.
- ❑ Documents are schema flexible.
- ❑ Documents have version number which are incremented with each update.

```
{  
  "_index": "blog",  
  "_id": "123456",  
  "_version": 2,  
  "_source": {  
    "title": "Introduction to Elasticsearch",  
    "author": "John Doe",  
    "publish_date": "2022-01-15",  
    "content": "Elasticsearch is a scalable  
search and analytics engine built on  
Apache Lucene. It offers real-time search  
capabilities, making it popular for log  
analytics, full-text search, and more. Its  
flexible data model and distributed  
architecture support powerful querying  
and analysis."  
  }  
}
```

What We Cover.

1. An **Introduction** to Elasticsearch.
2. Deep Dive into Elasticsearch's **Architecture**.
3. Elasticsearch in Action: **How Companies Utilize It**.
4. **Advanced Search Techniques** with Elasticsearch.
5. **Kibana** and Elasticsearch: The Perfect Combination.
6. **Conclusion**: Key Takeaways on Elasticsearch.

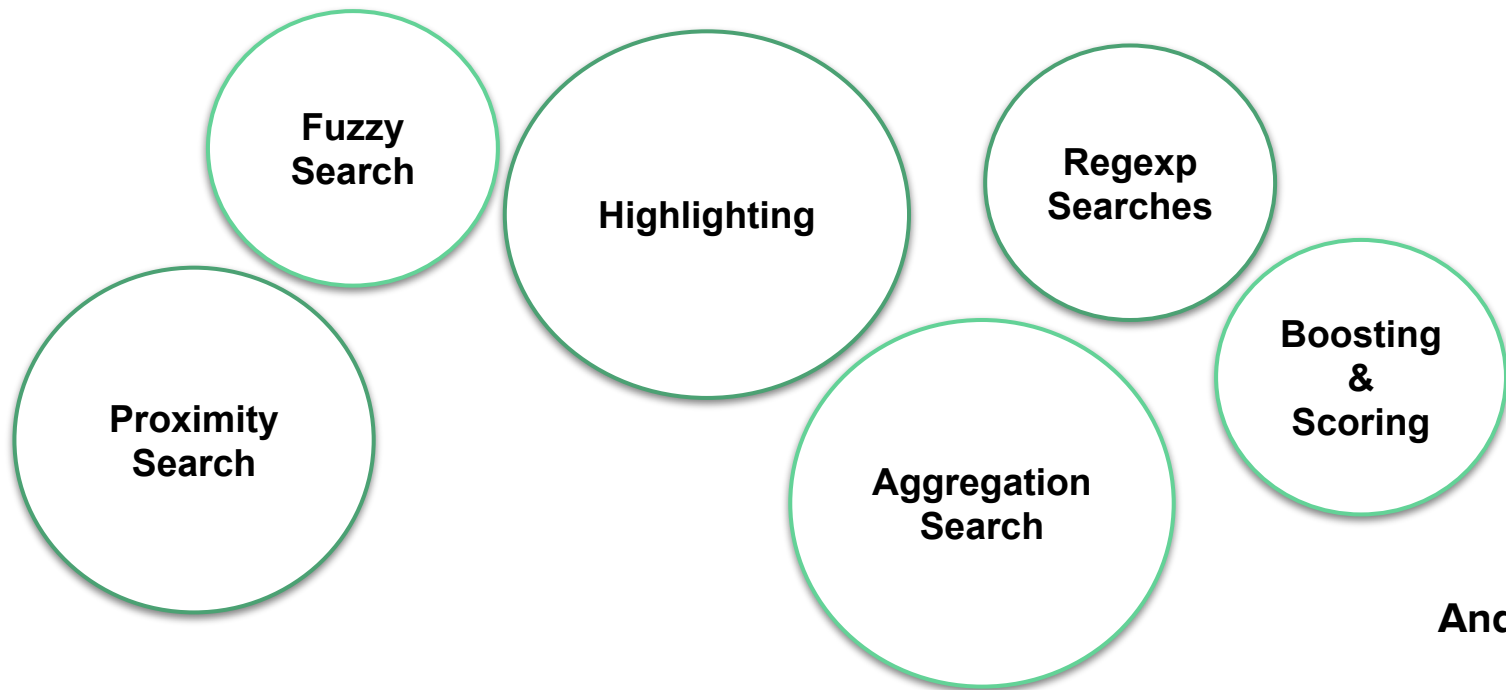
Elasticsearch in Action: **How Companies Utilize It.**

- ❑ Vodafone (EG) used it as a log analytics and reporting tool for their Website and App (personal experience).
- ❑ eBay's search infrastructure is powered by Elasticsearch help them deliver relevant results.
- ❑ NASA uses Elasticsearch to analyze and search through large volumes of telemetry data collected from satellites and space missions.
- ❑ Uber uses Elasticsearch to power real-time geospatial search.

What We Cover.

1. An **Introduction** to Elasticsearch.
2. Deep Dive into Elasticsearch's **Architecture**.
3. Elasticsearch in Action: **How Companies Utilize It**.
4. **Advanced Search Techniques** with Elasticsearch.
5. **Kibana** and Elasticsearch: The Perfect Combination.
6. **Conclusion**: Key Takeaways on Elasticsearch.

Advanced Search Techniques with Elasticsearch.




And many more...

What We Cover.

1. An **Introduction** to Elasticsearch.
2. Deep Dive into Elasticsearch's **Architecture**.
3. Elasticsearch in Action: **How Companies Utilize It**.
4. **Advanced Search Techniques** with Elasticsearch.
5. **Kibana** and Elasticsearch: The Perfect Combination.
6. **Conclusion**: Key Takeaways on Elasticsearch.

Kibana and Elasticsearch: The Perfect Combination.

Elastic Stack



What is Kibana?

- ❑ **Visualization** capabilities on top of indexed content.
- ❑ **Geospatial analysis** tool for creating and layering maps.
- ❑ **Development tools** for developers to interact with the Elasticsearch REST API.
- ❑ Users can create and manage **alerts** that notify them when real-time data meets certain conditions.

Kibana and Advanced Search: A **Demo.**

What We Cover.

1. An **Introduction** to Elasticsearch.
2. Deep Dive into Elasticsearch's **Architecture**.
3. Elasticsearch in Action: **How Companies Utilize It**.
4. **Kibana** and Elasticsearch: The Perfect Combination.
5. **Advanced Search Techniques** with Elasticsearch.
6. **Conclusion**: Key Takeaways on Elasticsearch.

Conclusion: Key Takeaways on Elasticsearch

- ❑ **Elasticsearch:** A potent, distributed search and analytics engine, underpinned by Apache Lucene.
- ❑ **Distributed Architecture:** Ensures high availability and resilience, with a structure composed of nodes, indices, shards, and replicas.
- ❑ **Elastic Stack Synergy:** Elasticsearch operates in concert with Kibana, Logstash, and Beats to provide a holistic data analysis solution.

Key Takeaways on Elasticsearch (Cont'd)

- ❑ **Advanced Search:** Employs sophisticated techniques such as fuzzy search, proximity matching, boosting, and aggregations for precise and intricate searches.
- ❑ **Customizable Scoring:** Allows tailoring of relevance scoring to cater to specific needs.
- ❑ **Versatility:** Used for a broad range of applications, including full-text search, data analysis, log, and event data management.
- ❑ **Scalability and Resilience:** Its ability to scale and remain resilient makes Elasticsearch a preferred choice among various organizations.



Elasticsearch: Going further

[Official Elasticsearch Documentation](#)
[Tutorialspoint's Elasticsearch tutorial](#)
[Tutorial series created by LisaHJung](#)

Thank you for your attention!